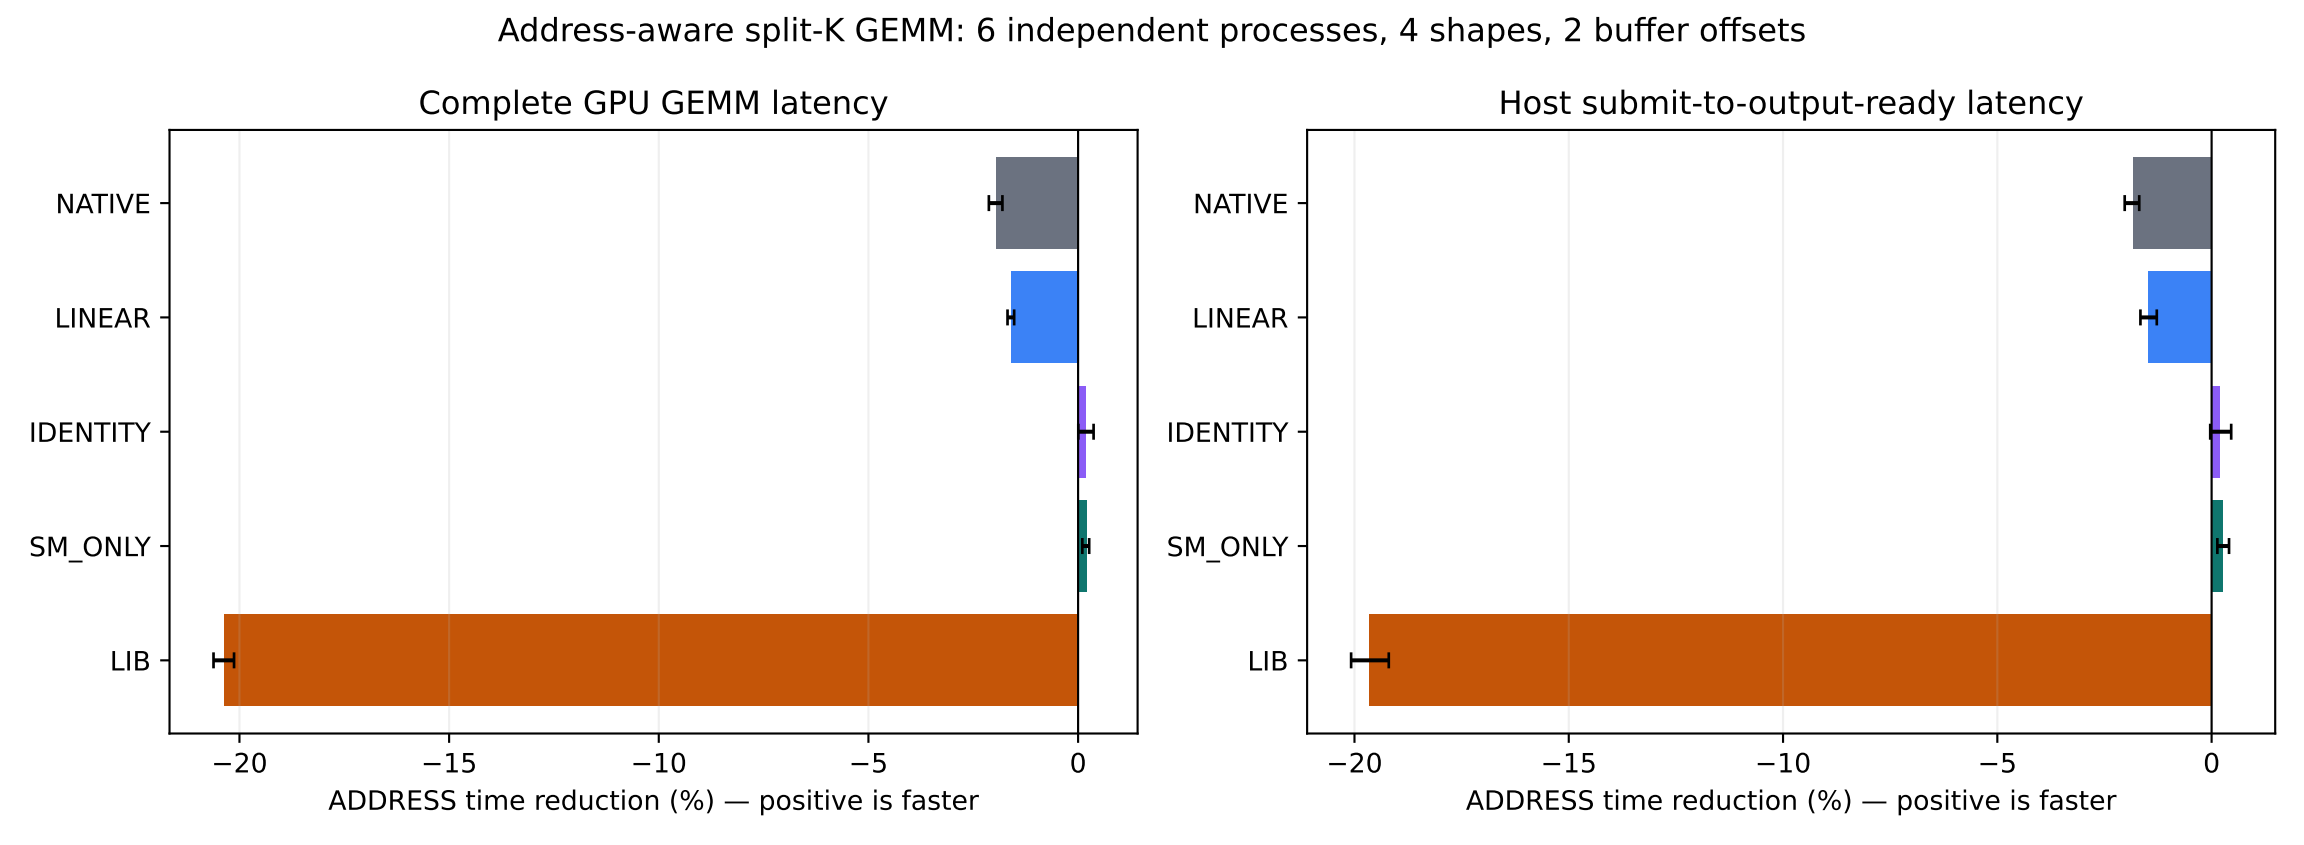
Address-aware split-K GEMM: 6 independent processes, 4 shapes, 2 buffer offsets
Complete GPU GEMM latency
Host submit-to-output-ready latency
NATIVE
LINEAR
IDENTITY
SM_ONLY
LIB
−20
−15
−10
−5
0
ADDRESS time reduction (%) — positive is faster