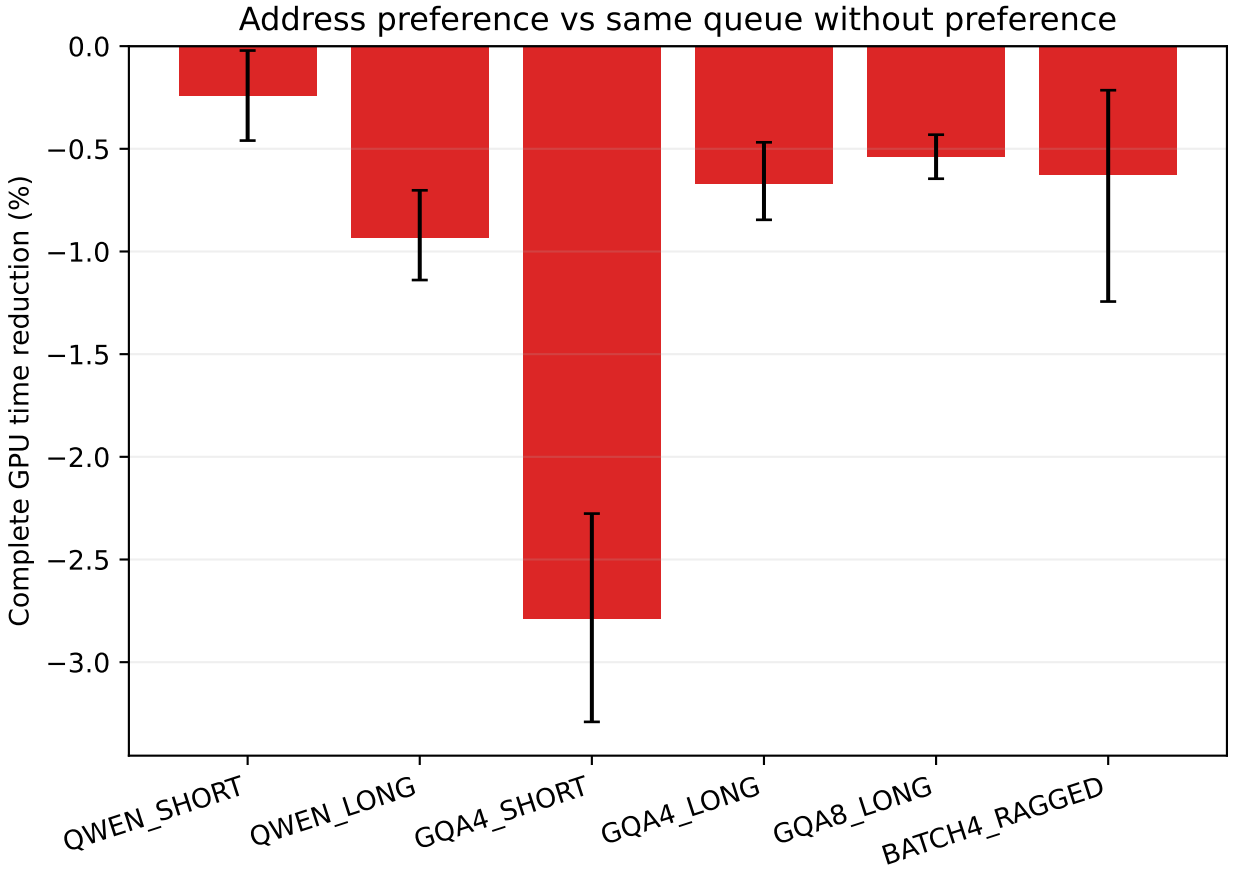
Address preference vs same queue without preference
Complete GPU time reduction (%)
0.0
−0.5
−1.0
−1.5
−2.0
−2.5
−3.0
QWEN_SHORT
QWEN_LONG
GQA4_SHORT
GQA4_LONG
GQA8_LONG
BATCH4_RAGGED

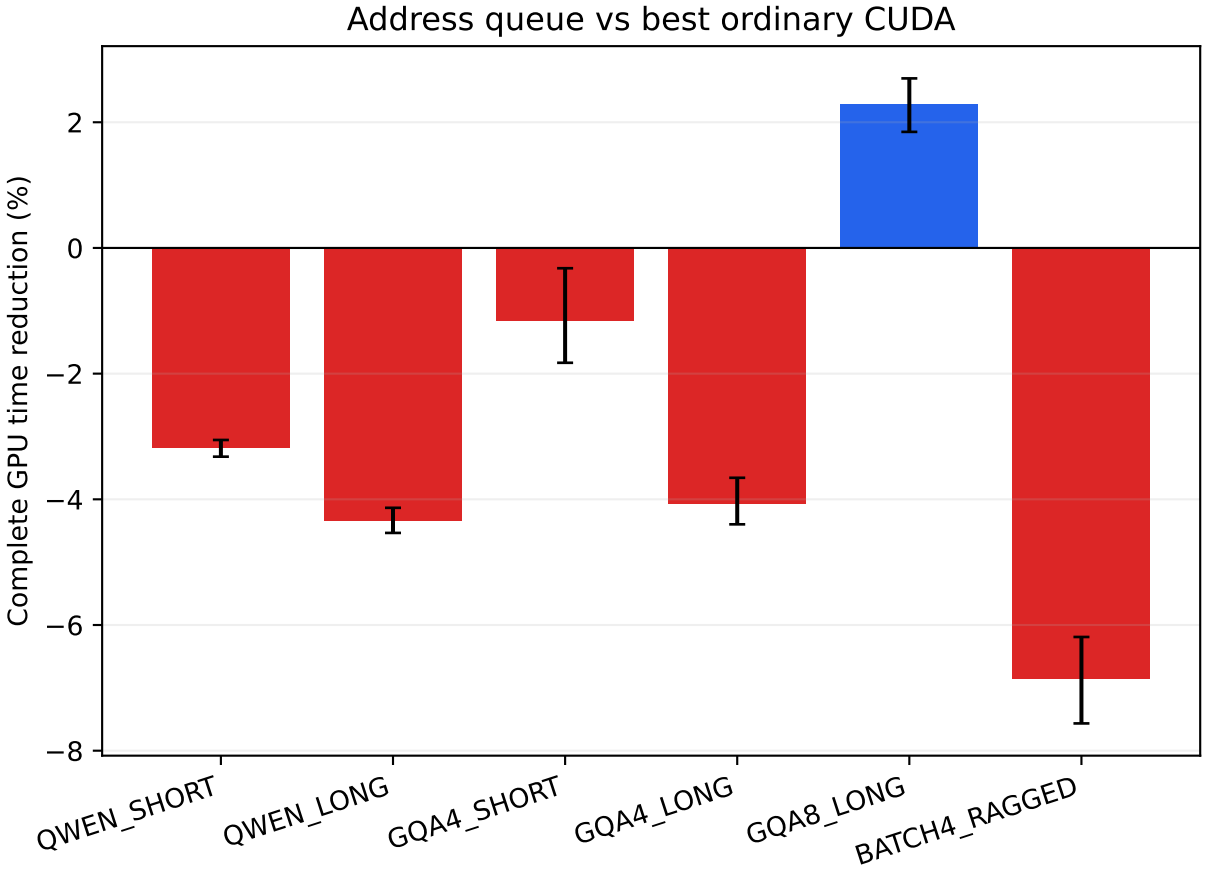
Address queue vs best ordinary CUDA
Complete GPU time reduction (%)
2
0
−2
−4
−6
−8
QWEN_SHORT
QWEN_LONG
GQA4_SHORT
GQA4_LONG
GQA8_LONG
BATCH4_RAGGED