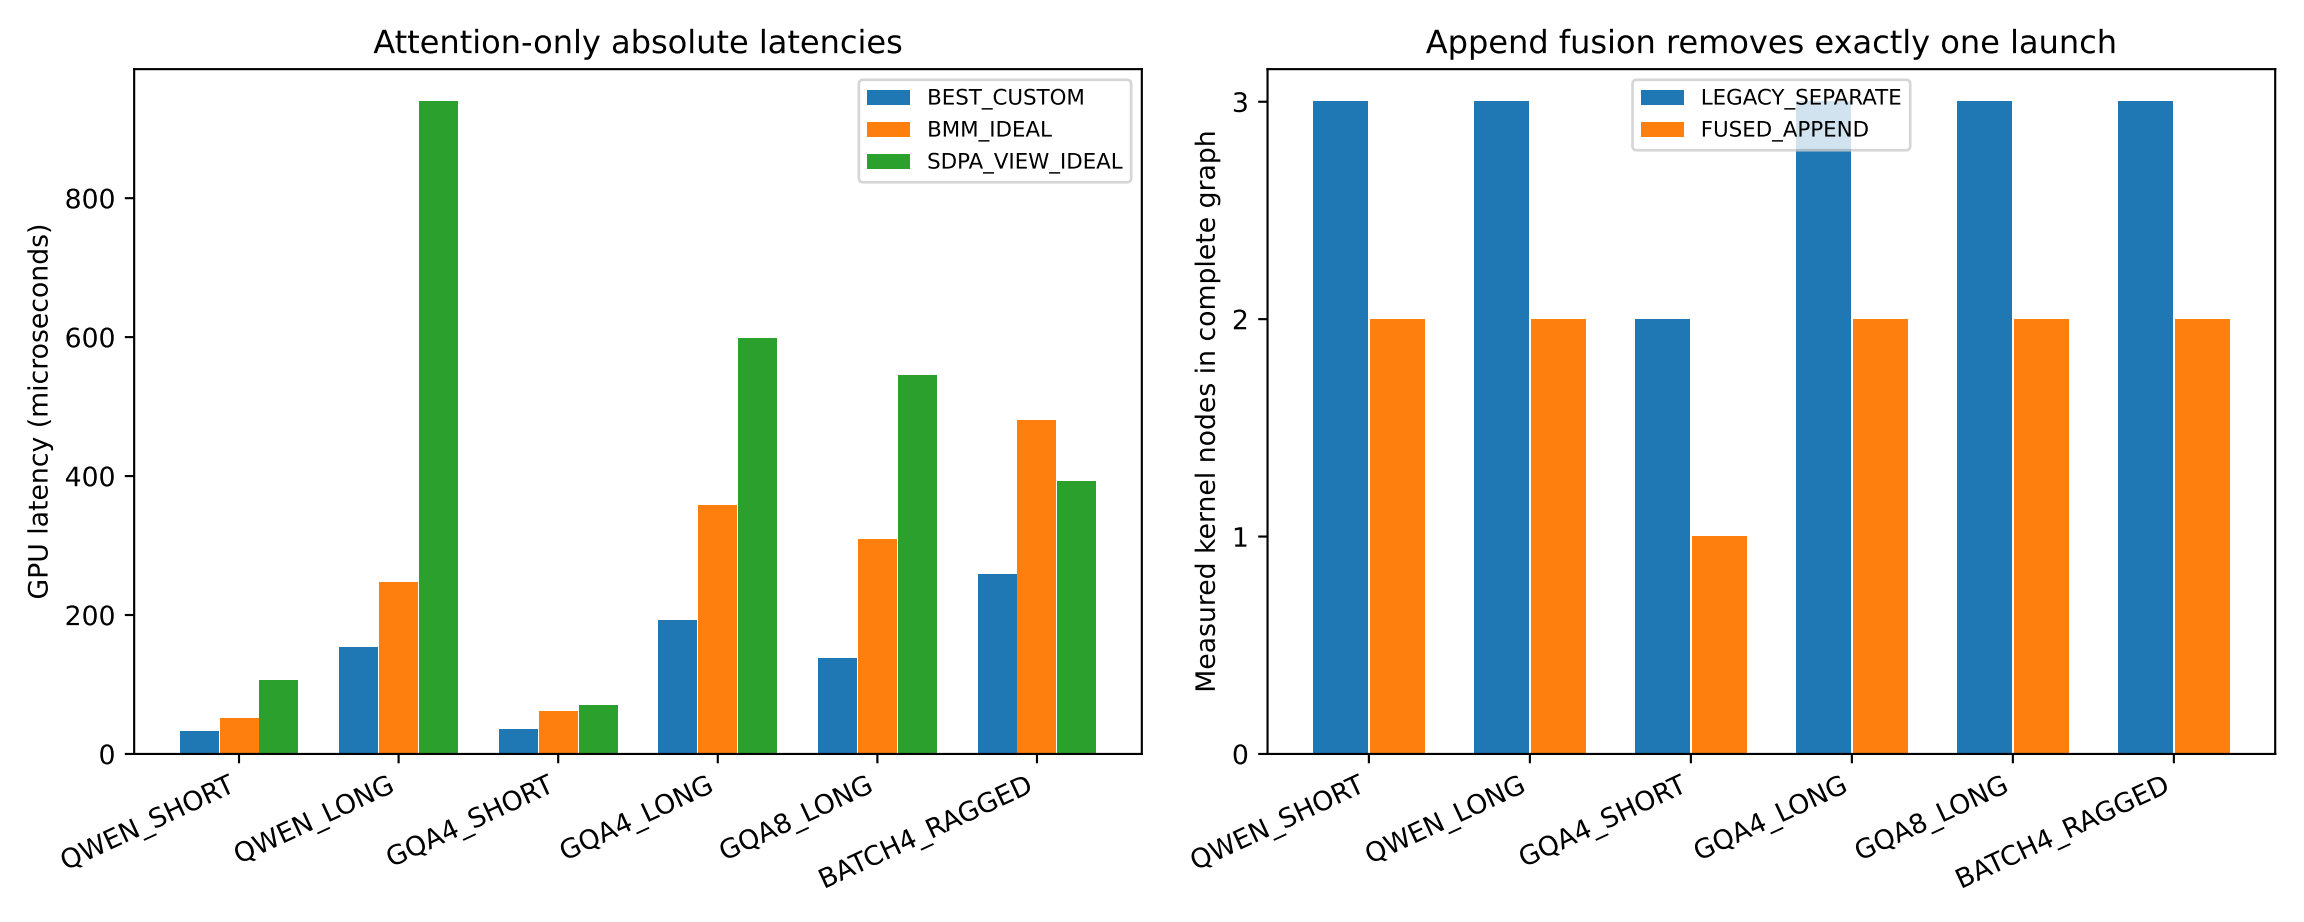
Attention-only absolute latencies
GPU latency (microseconds)
0
200
400
600
800
BEST_CUSTOM
BMM_IDEAL
SDPA_VIEW_IDEAL
QWEN_SHORT
QWEN_LONG
GQA4_SHORT
GQA4_LONG
GQA8_LONG
BATCH4_RAGGED
Append fusion removes exactly one launch
Measured kernel nodes in complete graph
0
1
2
3
LEGACY_SEPARATE
FUSED_APPEND
QWEN_SHORT
QWEN_LONG
GQA4_SHORT
GQA4_LONG
GQA8_LONG
BATCH4_RAGGED